

The Copy-Paste Complaint

Exercises on the CFPB Complaint Vault Dataset (18.2 million U.S. consumer complaints)

Name: _____ Class: _____ Date: _____

Level: upper secondary or first-year university. **Subjects:** statistics, economics, personal finance, civics, data literacy. **Tools:** Excel or Google Sheets (pivot tables or SUMIFS). No programming needed.

Before you start

The U.S. Consumer Financial Protection Bureau (CFPB) publishes every consumer complaint it receives about financial products. This dataset adds three things: a product taxonomy that is comparable across years, a measure of which complaint narratives are copy-paste templates, and company and bank rates. A complaint is an allegation, not a finding, and company responses are self-reported.

Files used (free to download from the Kibbo Consumer Data Library at getkibbo.com and from the Zenodo record, doi.org/10.5281/zenodo.23188842):

- complaints-by-year.csv (year, n)
- complaints-by-year-product-family.csv (year, product_family, n)
- complaints-by-year-product-family-response.csv (year, product_family, company_response_to_consumer, n)
- template-category-by-year-product-family-narratives-only.csv (year, product_family, template_category, n, pct)
- bank-complaint-rates.csv (company, year, complaint counts, FDIC deposits and account counts, and the rates built from them)
- top100-template-clusters.csv (extra exercise)

Read this first:

- 2026 is a partial year: complaints received from 1 January to 3 October. For narratives, the CFPB stopped publishing on 14 August 2026.
- “In progress” complaints have no final outcome yet. Exercises that use outcomes exclude them.
- A complaint is one person’s account. A company’s response category is reported by the company.

Exercises

Exercise 1 — Read, compare, and be careful with partial years (*basic*)

File: complaints-by-year.csv

1. How many complaints did the CFPB receive in 2024 and in 2025?
2. By what percentage did complaints change from 2024 to 2025?
3. The 2026 row is larger than the 2025 row. Does that mean 2026 had more complaints than 2025? Explain.
4. Calculate average complaints per day for 2025 (365 days) and for 2026 (1 January to 3 October = 276 days). Which year had more complaints per day, and by how much?

Exercise 2 — Proportions, and what “everything else” did (*basic*)

File: complaints-by-year-product-family.csv

1. Calculate credit reporting’s share of all complaints in 2015, 2025 and 2026.
2. Calculate how many complaints were **not** about credit reporting in 2015 and in 2025. How many times larger is the 2025 figure?
3. Money transfers & virtual currency: how many complaints in 2024 and in 2025? How many times larger is 2025?
4. Complete the sentence: “Credit reporting’s share rose from ___% to ___%, but complaints about everything else also _____.”

Exercise 3 — The mix effect (*intermediate*)

File: complaints-by-year-product-family-response.csv

Monetary relief rate = complaints closed “with monetary relief” ÷ all complaints with a final outcome (exclude “In progress”). Use 2023 and 2025.

1. Calculate the monetary relief rate for **all** complaints in 2023 and in 2025.
2. Calculate it again **excluding** credit reporting, and the rate for credit reporting alone.
3. Credit reporting’s share of complaints with a final outcome went from X% to Y%. Calculate both.
4. The overall rate fell much more than the rate excluding credit reporting. Explain why in two sentences, using the numbers.
5. (*Challenge*) If 2025 had the same product mix as 2023, but each group kept its own 2025 rate, what would the overall 2025 rate be? How much of the fall is explained by the change in mix, and how much by the change in rates within each group? Then redo the split using the 2025 mix with

2023 rates. Do you get the same split? Why or why not?

Exercise 4 — Template detection: measuring copy-paste (*intermediate*)

File: template-category-by-year-product-family-narratives-only.csv

Categories: **original** (unique text); **shared_template_strict** (same text filed from two or more different locations); **same_filer_repeat** (same text from one location, probably one person filing several times); **unknown_filer**. The file only covers complaints with a published narrative.

1. Calculate the share of narratives that were shared templates in 2015, 2020 and 2025.
2. In 2025, which product family had the highest shared-template share? Which had the lowest?
3. The 2026 share of shared templates is much lower than 2025. Give two reasons not to conclude that templates fell in 2026.
4. In 2025, what share of all narratives was about credit reporting? What share of all shared-template narratives was credit reporting?
5. Calculate the 2025 shared-template share **excluding** credit reporting. Compare it with the overall figure and explain what this tells you about where templates are concentrated.
6. Why does the dataset separate shared templates from same-filer repeats? What would happen to the template measure if it did not?

Exercise 5 — Denominators: three reasonable measures, three rankings (*advanced*)

File: bank-complaint-rates.csv. Use year = 2025 only.

Three ways to compare banks: **(A)** `complaints_per_billion_deposits_all` (all complaints per \$1 billion of deposits); **(B)** `bank_accounts_complaints_per_billion_deposits` (bank-account complaints only, per \$1 billion of deposits); **(C)** `bank_accounts_complaints_per_100k_accounts` (bank-account complaints per 100,000 deposit accounts of \$250,000 or less).

1. Keep only 2025 rows where `min_volume_flag` is FALSE and the per-100k rate is not blank. How many banks remain?
2. Rank those banks from highest to lowest on A, B and C. Which bank is first on each measure?
3. Find Citibank and Bank of America. What rank does each get on A, B and C? On which measures is Citibank's rate higher than Bank of America's, and by how many times?
4. The three measures give three different rankings. Which one best answers the question "which bank gets more complaints per customer?" Justify it, and list two weaknesses of that measure.
5. The rule that removes banks with fewer than 300 bank-account complaints is arbitrary. Raise the threshold to 500: which banks drop out, and who is first on C now? What does that tell you about

the reliability of rankings built on small counts?

Exercise 6 — What can we not conclude? (*advanced*)

Data: the tables below (Kibbo analysis of the complaint-level file, 2022–2025, comparing original and shared-template narratives within the same issue). Share of complaints closed with each type of relief, in %.

Closed with monetary relief

Issue	Original	Shared template
Bank accounts: Problem caused by your funds being low	18.6	0.3
Bank accounts: Managing an account	13.5	0.2
Money transfers: Other transaction problem	4.6	0.0

Closed with non-monetary relief (for example, correcting a credit report)

Issue	Original	Shared template
Credit reporting: Incorrect information on your report	35.1	51.2
Debt collection: Attempts to collect debt not owed	16.8	38.6
Debt collection: Written notification about debt	14.8	46.1

1. For each row, calculate the gap in percentage points between original and shared template. For the second table, also calculate the ratio (template ÷ original).
2. A student writes: “Using a template makes it less likely you get money and more likely you get a correction.” Rewrite the sentence so it is supported by the data and no more.
3. List at least three reasons the data cannot show that the template itself caused the difference (think: who uses templates, what problems they have, how companies record responses).
4. A company refunds a consumer but records the complaint as “Closed with explanation”. How would that appear in the tables, and why does it matter?
5. Does a shared template mean the complaint is false? Use the data and the definitions to answer.

Connect it to real life

The tables show that, for credit report errors, shared-template complaints more often end with a correction. If you found an error on your credit report, what three specific details would you add to any template before sending it, and why? Then use the Kibbo Credit Report Dispute Letter Generator (getkibbo.com) to draft your letter. If your problem is a debt you do not recognise, use the Kibbo Debt Collection Dispute Letter Generator instead.

Extra exercise — Read the biggest clusters (*advanced*)

File: `top100-template-clusters.csv` (the 100 largest groups of near-identical narratives; the file includes the first 300 characters of each representative text).

1. Sort by size. Which cluster is the largest, and how many narratives and distinct locations does it have?
2. Which cluster covers the most distinct locations? What company is it mostly about?
3. For how many of the 100 clusters is credit reporting the main product family (the first one listed in `top3_product_families`)?
4. Create a column: $\text{distinct locations} \div \text{size}$. Which cluster has the lowest ratio? Give two possible explanations, and say what other data you would need to choose between them.
5. Clusters were built within each calendar year (and 2025 within each half). Why can you not compare cluster sizes directly across all periods?

The texts in this file are mass-repeated texts, not individual accounts. Look for patterns, not for individuals.